

CSC320 — Introduction to Visual Computing, Spring 2016

Assignment 2: Image Inpainting

Posted: Tuesday, March 29, 2016

Due: 11:59pm, Friday, April 8, 2016

Late policy: 15% marks deduction per 24hrs, submission not accepted if > 4 days late

Combined grade of programming assignments: $\max\left(\frac{0.12A_1+0.13A_2}{0.25}, \frac{0.12A_1+0.25A_2}{0.37}\right)$

In this assignment you will implement and experiment with an image inpainting tool. The tool is based on the *Exemplar-Based Image Inpainting* technique by Criminisi *et al.* and will be discussed in class and in tutorials this week. Your specific task is to complete the technique's implementation in the starter code. The starter code is itself an extension of what you used for Assignment 1 Part B and is based on Kivy and OpenCV.

Goals: The goals of the assignment are to (1) get you familiar with reading and understanding a research paper and (partially) implementing the technique it describes; (2) learn how to implement basic operations such as computing image gradients and curve normals; and (3) learn how to assess your implementation's correctness and the overall technique's failure points.

Weighting of A2: Since we had only two assignments instead of three this semester, I decided to allow two weights for this assignment—either 12% (the originally-posted weight of A2) or 25% (the originally-posted weight of A3)—and use whichever yields the highest grade.

Bonus components: Given the brevity of the period you have to work on the assignment, I am not including any explicit bonus components. If, however, you feel that your implementation/extensions perform much better than the reference solution please send me an email to discuss this.

Important: You are advised to start *immediately* by reading the paper (see below). The next step is to run the reference solution as well as the starter code, and compare the differences in how they behave (*e.g.*, their choice of patches for each iteration). The basic structure of the starter code is identical to what you have seen already in Assignment 1 and you only need to understand a relatively small part of the code you are given. Once you “get the hang of it,” the programming component of the assignment should not be too hard as there is relatively little python coding to do. What will take most of the time is internalizing exactly what you have to do, and how.

Exemplar-Based Image Inpainting (100 Marks)

The technique is described in full detail in the following paper (included with your starter code and also available [here](#)):

A. Criminisi, P. Pérez and K. Toyama, “Region Filling and Object Removal by Exemplar-Based Image Inpainting,” *IEEE Transactions on Image Processing*, vol. 13, no. 9, 2004.

You should read Sections I and II of the paper right away to get a general idea of the principles behind the method. Section II, in particular, is very important because it introduces the notation used in the rest of the paper as well as the starter code. Section III describes the algorithm in detail, with pseudocode shown in Table I. The starter code implements exactly what is shown in Table I; the only thing left for you to implement is the term $D(\mathbf{p})$ in Eq. (1). Sections IV and V are not strictly necessary to read, but they do show many results that should give you more insight into how your implementation is supposed to behave.

Part 0. Starter code & the reference solution

Use the following sequence of commands on CDF to unpack the starter code:

```
> cd ~
> tar xvfz inpainting.tar.gz
> rm inpainting.tar.gz
```

Consult the file *320/A2/code/README_1st.txt* for details on how the code is structured and for guidelines about how to navigate it (the structure is very similar to A1 Part B). In addition to the starter code, I am providing a fully-featured reference solution in compiled, statically-linked binary format (OS X and CDF/Linux only). You should run this binary to see how your own implementation should behave, and to make sure that your implementation produces the correct output. That being said, you should not expect your implementation to produce *exactly* the same output as the reference solution as tiny differences in implementation might lead to slightly different results. This is not a concern, however, and the TAs will be looking at your code as well as its output to make sure what you are doing is reasonable.

320/A2/CHECKLIST.txt: Please read this form carefully. It includes information on the course's Academic Honesty Policy and contains details the distribution of marks in the assignment. You will need to complete this form prior to submission, and will be the first file markers look at when grading your assignment.

Part 1. Programming Component (80 Marks)

You need to complete the implementation of three functions detailed below. A skeleton of all three is included in file *320/A2/code/inpainting/compute.py*. This file is where your entire implementation will reside.

In addition to these functions, you will need to copy a few lines of code from your A1 implementation that are not provided in the starter code. This requires no effort other than verbatim line-by-line copy from your existing code. See *320/A2/code/README_1st.txt* for details.

Part 1.1. Computing Gradients: The *computeGradient()* function (30 Marks)

This function takes three input arguments: (1) a patch $\Psi_{\mathbf{p}}$ from the image being inpainted, represented as a member of the class *PSI* from file *code/inpainting/psi.py*; (2) a binary OpenCV image F indicating which pixels have already been filled; and (3) the color OpenCV image I being inpainted. The function returns the gradient with the largest magnitude within patch $\Psi_{\mathbf{p}}$:

$$\text{computeGradient}(\Psi_{\mathbf{p}}, F, I) := \nabla \tilde{I}_{\mathbf{q}^*} \quad (1)$$

$$\text{where } \mathbf{q}^* = \arg \max_{\substack{\mathbf{q} \in \Psi_{\mathbf{p}} \\ F(\mathbf{q}) > 0 \\ \nabla \tilde{I}_{\mathbf{q}} \text{ is valid}}} |\nabla \tilde{I}_{\mathbf{q}}| \quad (2)$$

and all gradients are computed on the grayscale version, \tilde{I} , of color image I .

If image I was a regular image with no “missing” pixels, implementing this function would be trivial with OpenCV: you would just need the OpenCV function that converts color images (or image patches) to grayscale, and the OpenCV function that computes the horizontal and vertical components of the gradient.

The complication here is that not all pixels \mathbf{q} in $\Psi_{\mathbf{p}}$ have been filled. As a result, applying those OpenCV functions will produce estimates of $\nabla \tilde{I}_{\mathbf{q}}$ that are incorrect/invalid for some pixels \mathbf{q} in that patch. Your main task in implementing `computeGradient()` will therefore be to find a way to ignore those pixels so that the max operation in Eq. (2) is not corrupted by these invalid estimates.

Efficiency considerations: You should pay attention to the efficiency of the code you write but you will not be penalized for correct, yet inefficient, solutions. *Hint:* The reference implementation is 10-12 lines of code and includes no explicit looping over pixels.

Part 1.2. Computing Curve Normals: The `computeNormal()` function (40 Marks)

Much like the previous function, this function also takes three input arguments and returns a 2D vector. The function's arguments are (1) a patch $\Psi_{\mathbf{p}}$ from the image being inpainted; (2) a binary OpenCV image F indicating which pixels have already been filled; and (3) a binary OpenCV image that is non-zero only for the pixels on the fill front $\delta\Omega$. The function assumes that the patch's center \mathbf{p} is on the fill front and returns the fill front's normal $\mathbf{n}_{\mathbf{p}}$ at that pixel.

You are free to use whatever method you wish for estimating the curve normal. This includes using finite differences between \mathbf{p} and its immediate neighbors on the fill front to estimate the tangent and normal at \mathbf{p} ; fitting a curve to the fill front in the neighborhood of \mathbf{p} and returning its normal at \mathbf{p} ; and using any built-in OpenCV functions you wish for parts of these computations.

Hint: Depending on how it is implemented this function may involve a couple dozen lines of code, but can potentially be much less (and, of course, much more).

Part 1.3. Computing Pixel Confidences: The `computeC()` function (10 Marks)

This function takes three input arguments and returns a scalar. The function's arguments are (1) a patch $\Psi_{\mathbf{p}}$ from the image being inpainted; (2) a binary OpenCV image F indicating which pixels have already been filled; and (3) an OpenCV image that records the confidence of all pixels in the inpainted image I . It returns the value of function $C(\mathbf{p})$ shown on page 4 of the paper, *i.e.*, it assumes that the confidence of unfilled pixels is zero and returns average confidence of all pixels in patch $\Psi_{\mathbf{p}}$.

Hint: The reference implementation is less than 5 lines of code and includes no explicit looping over pixels.

Part 2.1. Experimental evaluation (10 Marks)

Your task here is to put the inpainting method to the test by conducting your own experiments. Specifically, you need to do the following:

1. Run your inpainting implementation for 100 iterations on the test image pairs (*input-color.jpg*, *input-alpha.bmp*), and (*Kanizsa-triangle-tiny.png*, *Kanizsa-triangle-mask-tiny.png*). Save the partial results to directory *320/A2/report* using the following file naming convention: *X.inpainted.png*, *X.fillFront.png*, *X.confidence.png*, *X.filled.png* where *X* is either *input* or *Kanizsa*.
2. Capture two photos, *Source1* and *Source2*, with your own camera. Any camera will do—cellphone, point-and-shoot, action camera, etc. Be sure to reduce their size to something on the order of 300×200 pixels or less so that execution time is manageable.
3. Create binary masks, *Mask1* and *Mask2*, to mask out one or more elements in these photos. Use any tool you wish for this task.

4. **Important:** Your source photos and masks should *not* be arbitrary. You should choose them as follows: (a) the region(s) to be deleted should not be just constant-intensity regions, which are trivial to inpaint; (b) the pair *Source1,Mask1* should correspond to a “good” case for inpainting, *i.e.*, should be possible to find a patch radius that produces an inpainted image with (almost) no obvious seams or other visible artifacts; and (c) the pair *Source2,Mask2* should correspond to a “bad” case for inpainting, *i.e.*, it should not be possible to find a patch radius that produces an inpainted image free of very obvious artifacts.
5. Run the inpainting algorithm on these two input datasets. You are welcome to use the reference implementation for these experiments (*i.e.*, you don’t need to have your code fully functional in order to get started on this part of the assignment). Save the two sources, masks and inpainting results to directory *320/A2/report* using the file names *Source1.png*, *Mask1.png*, *Source1.inpainted.png* and *Source2.png*, *Mask2.png*, *Source2.inpainted.png*.

Part 2.2. Your PDF report (10 Marks)

Your report should include the following: (1) why you think the pair (*Source1,Mask1*) represents a good case for inpainting; (2) why you think the pair (*Source2,Mask2*) represents a bad case for inpainting; (3) discussion of any visible artifacts in the results from these two datasets (*i.e.*, what artifacts do you see and why you think they occurred).

Place your report in file *320/A2/report/report.pdf*. You may use any word processing tool to create it (Word, LaTeX, Powerpoint, html, *etc.*) but the report you turn in must be in *PDF format*.

What to turn in: Use the following sequence of commands on CDF to pack everything and submit:

```
> cd ~/CS320/
> tar cvfz assign2.tar.gz A2/CHECKLIST.txt A2/code A2/report
> submit -c csc320h -a assign2 assign2.tar.gz
```

Working on non-CDF machines: You are welcome to work on the assignment on a non-CDF machine. That being said, your implementation must run on the Linux CDF machines in order to be marked by the TAs. You should therefore test your code on CDF frequently, to make sure it has no CDF-specific issues.

All required packages have already been installed on CDF/Linux. Here are the steps I followed to install them on OS X 10.10+ :

1. Install NumPy and related packages for scientific computations:
<http://stronginference.com/ScipySuperpack/>
2. Install OpenCV 3.1:
<http://blogs.wcode.org/2014/10/howto-install-build-and-use-opencv-macosx-10-10/>
3. Install Kivy using homebrew and pip (needed for Part B only):
<https://kivy.org/docs/installation/installation-osx.html>

Freely-available resources on NumPy, OpenCV, Computer Vision Programming, etc:

1. Short NumPy tutorial by Olessia Karpova (CSC420, 2014):
https://github.com/olessia/tutorials/blob/master/numpy_tutorial.ipynb

2. Jan Erik Solem, *Programming Computer Vision with Python*, O'Reilly Media, 2012 (preprint):
<http://programmingcomputervision.com/> (look at Chapters 1 and 10)
3. *OpenCV-Python documentation* (Intro to OpenCV, Core Operations, Image Processing in OpenCV):
http://docs.opencv.org/3.1.0/d6/d00/tutorial_py_root.html#gsc.tab=0
4. *Matplotlib User Guide* (especially Chapter 5.5):
<http://matplotlib.org/1.5.1/Matplotlib.pdf>
5. *NumPy Reference Guide* (especially Chapters 1.4, 1.5, 3.17.1-3.17.5):
<http://docs.scipy.org/doc/numpy/numpy-ref-1.10.1.pdf>